The Representational Regress: Machine-Checked Topology, Diagonal Arguments, and Boundary Obstructions for Representational Systems

Nova Spivack
Independent Researcher

2026

Abstract

Punchline. A representational system—any structure that represents itself by iterating a fixed construction—faces a *topological* barrier, not merely a computational one. We prove, in machine-checked Lean 4 with zero `sorry` and no project-specific axioms, that the barrier cannot be removed by scaling, by deeper self-models, by partial self-interpreters, by cross-object encoding, by restricting to decidable predicates, or by any continuous deformation within a fixed topological type. The obstacle is **wrong type**, not **inexhaustible queue**: it is a finite-data invariant (a homeomorphism obstruction between boundary-incompatible spaces) that no amount of computation, iteration, or architectural extension can change.

Technical summary. The formalization treats *representational systems*—categories equipped with a non-trivial iterated self-map—and an intrinsic planar-chart predicate `ChartableR2` on compact topological models.

(i) **Categorical strand.** A representational system yields infinitely many formally distinct meta-levels; Lawvere's fixed-point theorem in `Type` and in `MonoidalClosed(Type)` is complemented by an inductive `OntologicalSlot` witness that separates morphism data from global-element data.

(ii) **Topological strand.** A general lemma states: any two spaces whose `ChartableR2`-co-characterized boundaries are subtype-incompatible are not homeomorphic. The closed cylinder and Möbius strip are a fully discharged compact witness.

(iii) **No-escape closures.** Six adversarial retreat routes are formally blocked: infinite regress, finite-order iterate collapse, cross-object section–retraction readout (both sides), partial (non-surjective) self-models, decidable/Bool-valued self-reference, and the original Lawvere enumerator obstruction. Each is a proved theorem, not an assumption.

(iv) **Bridges.** The library connects the obstruction to consciousness (under a conditional philosophical premise P2), symbol grounding, AI self-interpretation, and hypothetical observer–readout topologies—making explicit which domain-specific premises must be supplied before the barrier applies.

Consequence. In every neighboring debate—consciousness, symbolic AI, semantic grounding, measurement chains—admissible solutions share a common shape: *structural change* to the topological type of the substrate, not quantitative extension within a fixed type. The formalism makes this precise and machine-checkable.

1 Introduction

1.1 The Problem

When a system represents itself—by iterating a fixed construction, by encoding its own states, by interpreting its own symbols—a familiar worry arises: infinite regress, or diagonal limitations, or both. But those concerns are usually framed *computationally*: the regress does not terminate, the self-interpreter cannot be total, the Gödel sentence is undecidable. The implicit hope is that with enough resources—more compute, richer architectures, deeper self-models—the obstacle might shrink or be routed around.

This paper shows, formally and machine-checkably, that the obstacle is not computational at all. It is **topological**: a structural mismatch between the boundary organization of two types of space, which no amount of iteration, scaling, or continuous deformation within a fixed type can remove. And we close, one by one, every standard adversarial route by which one might hope to circumvent it.

The technical development stays in categorical and topological language (Lean 4, checked against Mathlib, zero `sorry`, no project-specific axioms); §7.3–§7.7 relate the results to consciousness, symbolic grounding, diagonal arguments, AI scaling, and observer–apparatus topologies.

1.2 Summary of Results

The core result is a **universal topological obstruction** for representational systems—and a proof that the six most natural adversarial attempts to evade it all fail.

The obstruction. A general lemma (Theorem 4.4) states: if two spaces admit subsets characterized as exactly the non-`ChartableR2` points, and if those boundary *subspaces* are not homeomorphic, then the ambient spaces are not homeomorphic. This is a *finite-data invariant*: two spaces are either homeomorphic or not, regardless of how many modeling steps one enumerates. The closed cylinder and Möbius strip are a fully discharged compact witness (**M-FINAL**). Any future domain whose state space matches the same boundary pattern imports the same conclusion; no separate bridge argument is required.

Why this is not "just" infinite regress. The categorical strand (§2) does yield infinitely many distinct meta-levels under the bundled hypothesis `iter_injective`, with Lawvere's wall and `OntologicalSlot` preventing collapse of the endo/obj distinction. But infinity is not the deepest point. The topological obstruction is independent: it says the barrier is **wrong type**, not **inexhaustible queue**. No amount of compute, no deeper self-model, no richer architecture changes the boundary pattern of a topological type.

No escape. Six adversarial retreat routes are formally closed (§7.6):

(i) Infinite injective tower ⇒ full regress + Lawvere wall.
(ii) Finite-order iterate collapse ⇒ bounded depth, no injective tower.
(iii) Cross-object readout via section–retraction ⇒ idempotent on host, identity on source; no injective tower on *either* side.
(iv) Partial (non-surjective) self-model ⇒ diagonal blind spot—the very information needed for self-capture is the information it cannot represent.
(v) Decidable / Bool-valued self-reference ⇒ Lawvere obstruction persists in the finite, two-valued setting.

(vi) Different topological model ⇒ the general lemma applies to *any* pair with incompatible `ChartableR2` boundaries, not just cylinder/Möbius.

Together these leave an adversary with a clean dichotomy: either an infinite injective regress with the Lawvere wall intact, or a finite loop too shallow for full self-capture—and in either case the topological barrier stands.

Bridges and conditional readings. Cross-domain structures (`SymbolSystem`, `LawvereRegressBridge`, bundled `ChartableR2BoundaryModel` instances, `QuantumObserverChainHypothesis`) make the obstruction's *hypothesis shape* explicit for grounding, diagonal, consciousness (under premise P2), and observer–readout readings. An extended master theorem adjoins the topological results to the categorical core in one checkable bundle.

Utility. For any application the workflow is: (1) encode the system's state space in sufficient topological detail to verify the boundary biconditionals, (2) compare boundary subspaces up to homeomorphism, (3) invoke the general lemma or a bundled model. If the premises hold, the obstruction applies; if they fail, the obstacle was never licensed. That conditional structure is what keeps the mathematics honest while delivering a *portable* gate.

1.3 Contributions and Novelty Claims

Individual ingredients are classical: Lawvere's diagonal (1969), textbook topology of the cylinder and Möbius strip, Cantorian regress patterns. The contribution is a **unified, adversarially checked, machine-verified** landscape in which all three coexist in one artifact, so that "compute harder" is not confusable with "change homeomorphism type." Specific novel elements:

(1) `OntologicalSlot` — inductive tagging of morphism vs. global-element slots; makes the Lawvere wall explicit inside the type discipline. [[Novel]]

(2) `ChartableR2` **spine + seam–equator charts** — intrinsic $\mathbb{R}^2$ chartability as a definable boundary sensor; constructive charts through the Möbius quotient seam including the equator via piecewise unfolding. [[Novel] technique]

(3) **General boundary-contrast obstruction** — a single reusable, machine-checked lemma: any two spaces with incompatible `ChartableR2`-co-characterized boundaries are not homeomorphic. Wired to a concrete compact witness and to the categorical bundle. To our knowledge, *not* standard in the philosophy-of-mind or AI-safety literatures. [[Novel]]

(4) **Cross-domain bridge structures** — `SymbolSystem`, `LawvereRegressBridge`, bundled `ChartableR2BoundaryModel` data, `QuantumObserverChainHypothesis`: these make the obstruction's hypothesis shape explicit per domain. [[Novel]]

(5) **No-escape horn closures** — six formally blocked adversarial escape routes (iterate collision, section–retraction on both sides, partial-model diagonal exclusion, decidable Lawvere, wall persistence at every depth). The integrated closure package is, to our knowledge, new. [[Novel]]

1.4 Paper Organization

Section 2: representational systems, regress, Lawvere, `OntologicalSlot`, iterate dichotomy, core master theorem. Sections 3–6: cylinder and Möbius models, `ChartableR2`, generic lemma, seam

charts, M-FINAL. Section 7: extended master, instantiations (consciousness, grounding, diagonal arguments, QM schematic, LLM scaling), no-escape horn analysis (§7.6), admissibility (§7.7). Sections 8–10: limitations, related work, conclusion. Appendices: Lean module map, key theorem index, reproduction.

2 The Representational Regress

2.1 The `RepresentationalSystem` Structure

Definition 2.1 (`Basic.lean`: `RepresentationalSystem`).
A representational system *is a tuple* $(C, A, \text{represent}, \text{iter_injective})$ *where:*

- C *is a category,*
- A *is an object of* C*,*
- $\text{represent} : A \to A$ *is an endomorphism,*
- iter_injective*: for all* $n \neq m$*,* $\text{represent}^n \neq \text{represent}^m$ *in* $\text{End}(A)$*.*

2.2 Non-Termination and Unboundedness

Theorem 2.2 ([Known math] `Regress.lean`). *For any representational system* R*:*

(a) (`regress_no_termination`) Distinct indices $n \neq m$ *yield distinct meta-level arrows.*

(b) (`regress_iterates_unbounded`) For every bound B*, there exists a level* $L > B$ *with a distinct iterate.*

(c) (`regress_is_infinite`, `meta_range_infinite`) The set of meta-levels is infinite.

Remark 2.3 (What if `iter_injective` fails?). *The hypothesis* iter_injective *is* bundled*, not derivable from a bare endomorphism. An adversary might weaken it: suppose* $\text{represent}^n = \text{represent}^m$ *for some* $n \neq m$*.* `NoEscapeClosure` *formalizes the consequence: let* $e := \text{End.of}(\text{represent})$*; if* $e^{n+d} = e^n$ *with* $d > 0$ *then every power* e^k *equals some* e^i *with* $i < n + d$ *(`IterateClosure.End.exists_lt_of_pow_collision`). `IterateClosure.End.pow_iterate_dichotomy` records the* ***dichotomy****: for every monoid element* e*, either (i) distinct natural-number powers are all distinct (an infinite injective* $\mathbb{N}$*-tower), or (ii) the set of* distinct *powers is bounded by a finite index. In case (ii) the system has only finitely many representational layers—it loops. Either route blocks "just enough" self-representation: one gives an infinite regress; the other gives a finite loop too impoverished to host injective iterate depth.*

2.3 The Lawvere Wall

Theorem 2.4 ([Known math] `Lawvere.lean`). *(`lawvere_fixed_point_Type`) If* $s : A \to (A \to B)$ *is surjective, then every* $f : B \to B$ *has a fixed point.*

Theorem 2.5 ([Known math] `LawvereCCCType.lean`). *(`lawvere_fixed_point_MonoidalClosedType`) The same conclusion holds when the universal-enumerator hypothesis is recast as surjectivity of* $\text{curry}(\text{lawvereBinary } s)$ *in* $\text{MonoidalClosed}(\text{Type } u)$*.*

2.4 The `OntologicalSlot` Construction

Definition 2.6 ([Novel] `Basic.lean`: `OntologicalSlot`). *For a category C with terminal object and an object A:*

$$\text{OntologicalSlot}(C, A) \;:=\; \text{endo}(f : A \to A) \;\mid\; \text{obj}(g : \top_C \to A)$$

an inductive type with two constructors, disjoint by definition.

Theorem 2.7 ([Novel] `FixedPoints.lean`). *(`lawvere_wall_is_not_dissolution`) For any representational system R and any global element* $\text{fp} : \top_C \to A$,

$$\text{OntologicalSlot.endo}(R.\text{represent}) \;\neq\; \text{OntologicalSlot.obj}(\text{fp}).$$

2.5 The Core Master Theorem

Theorem 2.8 (`Main.lean`). *(`representational_regress_master`) For any representational system R:*

$$\text{representational_regress_master_claim}(R)$$

i.e., the conjunction of unbounded regress, Lawvere wall non-dissolution, and T_2 invariance under homeomorphism.

3 Topological Models: Cylinder and Möbius Strip

3.1 The Closed Cylinder

Definition 3.1 (`CylinderBoundary.lean`). ClosedCylinder $:= S^1 \times [0, 1]$ *with the product topology and manifold structure* $\mathfrak{R}_1 \times \mathfrak{R}_1^{\partial}$.

Theorem 3.2 ([Known math] `CylinderBoundary.lean`). *The manifold boundary of* ClosedCylinder *decomposes as:*

$$\partial(\text{ClosedCylinder}) = F_{\text{bot}} \sqcup F_{\text{top}}$$

where $F_{\text{bot}} = S^1 \times \{0\}$ and $F_{\text{top}} = S^1 \times \{1\}$ are each connected, and their union is disconnected.

3.2 The Möbius Strip

Definition 3.3 (`CylinderMobius.lean`). *The* Möbius equivalence $\sim$ *on* $[0, 1]^2$ *is the equivalence relation generated by* $(0, t) \sim (1, 1-t)$ *for* $t \in [0, 1]$, *extended by reflexivity, symmetry, and transitivity. The* Möbius strip *is* MobiusStrip $:= [0, 1]^2/\!\sim$.

Theorem 3.4 (`CylinderMobius.lean`).

(a) *(`mobiusRelOGraph_isClosed`) The graph of $\sim$ is closed in $[0, 1]^2 \times [0, 1]^2$.*

(b) *(`CompactSpace MobiusStrip`)* MobiusStrip *is compact (quotient of compact).*

(c) *(`instT2SpaceMobiusStrip`)* MobiusStrip *is Hausdorff (closed graph + compact domain).*

(d) *(`isConnected_mobiusStripBoundary`) The boundary band $\partial_{\text{model}} := \pi(\{t = 0\} \cup \{t = 1\})$ is connected (one component).*

3.3 Boundary Subspace Contrast

Theorem 3.5 (`MobiusCylinderBoundaryContrast.lean`).
(`mobiusStripBoundary_not_homeomorphic_closedCylinderBoundaryUnion`)

$$\mathrm{IsEmpty}(\mathrm{Subtype}(\partial_{\mathrm{Mob}}) \simeq_t \mathrm{Subtype}(\partial_{\mathrm{Cyl}}))\,.$$

The Möbius boundary band (connected) and the cylinder boundary union (disconnected) are not homeomorphic as subspaces.

4 The `ChartableR2` Framework

4.1 Definition and Basic Properties

Definition 4.1 (`ChartableR2Neighbor.lean`). *A point x of a topological space X is* ChartableR2 *if there exists an open set $U \ni x$ such that $U \cong_t \mathbb{R}^2$.*

Theorem 4.2 (`ChartableR2Neighbor.lean`).

(a) (`isOpen_setOf_chartableR2`) The set of `ChartableR2` points is open.

(b) (`isClosed_setOf_not_chartableR2`) Its complement is closed.

(c) (`not_chartableR2_of_tendsto_seq_atTop`) Sequential stability: limits of non-chartable points are non-chartable.

4.2 Invariance Under Homeomorphism

Theorem 4.3 (`ChartableR2Bridge.lean`).
(`Homeomorph.chartableR2_iff`)
If $h : X \simeq_t Y$ is a homeomorphism, then $\mathrm{ChartableR2}(x) \iff \mathrm{ChartableR2}(h(x))$.

4.3 The Generic Boundary-Contrast Obstruction

Theorem 4.4 ([Novel] `ChartableR2Bridge.lean`).
(`not_homeomorphic_of_chartableR2_boundary_contrast`)
Let X and Y be topological spaces with predicates $b_X \subseteq X$ and $b_Y \subseteq Y$ such that:

(i) $\forall x,\ x \in b_X \iff \neg\,\mathrm{ChartableR2}(x)$,

(ii) $\forall y,\ y \in b_Y \iff \neg\,\mathrm{ChartableR2}(y)$,

(iii) $\mathrm{IsEmpty}(\mathrm{Subtype}(b_X) \simeq_t \mathrm{Subtype}(b_Y))$.

Then $\mathrm{IsEmpty}(X \simeq_t Y)$.

Proof sketch. Any homeomorphism $e : X \simeq_t Y$ preserves ChartableR2 (Theorem 4.3), hence maps b_X to b_Y pointwise. By e.subtype, we get a homeomorphism $\mathrm{Subtype}(b_X) \simeq_t \mathrm{Subtype}(b_Y)$, contradicting (iii). □

Theorem 4.4 subsumes the non-homeomorphism strategy of §6: any other pair of spaces satisfying the same three hypotheses is also not homeomorphic.

4.4 Half-Space Obstructions

Theorem 4.5 (`HalfPlaneVsPlane.lean`).
(`euclideanHalfSpace2_not_homeomorphic_euclidean2`)
IsEmpty$(H^2 \simeq_t \mathbb{R}^2)$. *Proof via punctured half-plane (star-convex $\Rightarrow$ simply connected) vs. punctured plane (not simply connected, from* exp *covering).*

Theorem 4.6 (`HalfSpaceNeighborVsPlane.lean`).
(`isEmpty_homeomorph_subtype_openH2_zero_nbhd_subtype_openPlane`)
For open $V \subseteq H^2$ with $0 \in V$ and open $O \subseteq \mathbb{R}^2$, there is no homeomorphism $V \cong_t O$.

4.5 Cylinder C4: Boundary Biconditional

Theorem 4.7 (`CylinderChartableBoundary.lean`).
(`closedCylinder_boundaryUnion_iff_not_chartableR2`)
For all $x \in$ ClosedCylinder*:*

$$x \in \partial_{\mathrm{Cyl}} \iff \neg\,\mathrm{ChartableR2}(x).$$

5 Interior Charts on the Möbius Strip: The Seam Problem

5.1 Why the Seam Is Hard

Interior-height points of the fundamental domain fall into three regimes: strict interior $0 < x < 1$, where the quotient is injective on small boxes; the vertical seam $x \in \{0, 1\}$ with $t \neq \frac{1}{2}$, where trigonometric coordinates can separate fibers when the height window avoids the equator; and the seam at $t = \frac{1}{2}$, where $(0, \frac{1}{2}) \sim (1, \frac{1}{2})$ and a trigonometric chart degenerates. ChartableR2 at the equator quotient class is obtained by a piecewise Euclidean unfolding map (see below).

5.2 Open-Cell Charts

Theorem 5.1 (`MobiusChartableBoundary.lean`).
(`chartableR2_mobiusQuotientMk_of_Ioo_square`)
If $p \in (0, 1)^2$ (strict interior of the fundamental domain), then ChartableR2$([\![p]\!])$.

5.3 Seam Trig Charts (Away from Equator)

Definition 5.2 (`MobiusSeamChart.lean`, `MobiusSeamChartable.lean`). *The* saturated seam patch *$S(t_0, \varepsilon, \delta)$ is the union of the left and right ε-thickened seam strips at height window $(t_0 - \delta, t_0 + \delta)$, saturated under $\sim$.*

Theorem 5.3 (`MobiusSeamTrigInject.lean`).
(`onPatch_mobiusFDTrigCoord_eq_iff_mobiusRel0`)
When $\delta < |t_0 - \frac{1}{2}|$, the trigonometric FD coordinate map mobiusFDTrigCoord *satisfies:*
mobiusFDTrigCoord(p) = mobiusFDTrigCoord(q) *on the saturated patch if and only if $p \sim q$.*

Theorem 5.4 (`MobiusSeamChartableR2.lean`).
(`chartableR2_mobiusQuotientMk_of_mem_mobiusSeamSaturatedPatch`)
For $t_0 \neq \frac{1}{2}$ and appropriate ε, δ, every point
in $\pi(S(t_0, \varepsilon, \delta))$ is ChartableR2*.*

5.4 The Equator Chart: Piecewise Unfolding

Definition 5.5 (`MobiusSeamChartableR2.lean`).
The fundamental-domain unfolding map
$\mathrm{unfoldFD} : [0,1]^2 \to \mathbb{R}^2$ *is:*

$$\mathrm{unfoldFD}(x,t) := \begin{cases} \mathrm{mobiusPlaneCoord}(x,t) & \text{if } x \le \frac{1}{2}, \\ \mathrm{mobiusPlaneTrigSeamPartnerSub} & \text{if } x > \frac{1}{2}. \\ \quad \big(\mathrm{mobiusPlaneCoord}(x,t)\big) & \end{cases}$$

Lemma 5.6 (`MobiusSeamChartableR2.lean`). unfoldFD *is* $\sim$*-invariant: if* $p \sim q$ *then* $\mathrm{unfoldFD}(p) = \mathrm{unfoldFD}(q)$.

Theorem 5.7 (`MobiusSeamChartableR2.lean`).
(`chartableR2_mobiusQuotientMk_of_left_seam_half`)
$\mathrm{ChartableR2}([\![(0, \frac{1}{2})]\!])$.

5.5 Assembly: All Interior Points Are ChartableR2

Theorem 5.8 (`MobiusSeamChartableR2.lean`).
(`chartableR2_mobiusQuotientMk_of_interior_height`)
For every $p \in [0,1]^2$ *with* $0 < p_2 < 1$*:*

$$\mathrm{ChartableR2}([\![p]\!]).$$

5.6 Möbius C4: Boundary Biconditional

Theorem 5.9 (`MobiusChartableBoundary.lean`).
(`mobiusStripBoundary_iff_not_chartableR2_of_forall_off_edge_chartable`)
For all $x \in \mathrm{MobiusStrip}$*:*

$$x \in \partial_{\mathrm{Mob}} \iff \neg\, \mathrm{ChartableR2}(x).$$

6 M-FINAL: The Möbius Strip Is Not Homeomorphic to the Closed Cylinder

Theorem 6.1 (`MobiusSeamChartableR2.lean`, `ChartableR2Bridge.lean`).
(`mobiusStrip_not_homeomorphic_closedCylinder`)

$$\mathrm{IsEmpty}(\mathrm{MobiusStrip} \simeq_t \mathrm{ClosedCylinder})\,.$$

Proof sketch. Instantiate Theorem 4.4 with $b_X = \partial_{\mathrm{Mob}}$, $b_Y = \partial_{\mathrm{Cyl}}$, the biconditionals (Theorems 4.7 and 5.9), and the boundary incompatibility (Theorem 3.5). □

7 The Extended Master Theorem and Instantiations

7.1 Formal Statement

Theorem 7.1 (`Main.lean`).
(`representational_regress_master_extended`)

For any representational system R:

$$\text{RepresentationalRegressMasterExtended}(R)$$

where this abbreviation expands to:

(i) representational_regress_master_claim(R)
(regress + wall + T_2 transport),

(ii) IsEmpty($H^1 \simeq_t \mathbb{R}^1$) *(1D half-space vs. line;* `representational_regress_topology_halfLineModel`*),*

(iii) IsEmpty(MobiusStrip $\simeq_t$ ClosedCylinder) *(M-FINAL).*

This *extends* `representational_regress_master`; it does not replace it.

7.2 Bundled boundary models

Definition 7.2 (`ChartableR2Neighbor.lean`)**.** *A* `ChartableR2` boundary model *on X is a subset $b_X \subseteq X$ together with a proof that $\forall x,\ x \in b_X \iff \neg\,\text{ChartableR2}(x)$.*

- The type `ChartableR2BoundaryModel` packages such data; the lemma `ChartableR2BoundaryModel.not_homeomorphic` restates Theorem 4.4 in bundled form.
- The development `ChartableR2ConcreteBoundaryModels` supplies `mobiusStripChartableR2BoundaryModel` and `closedCylinderChartableR2BoundaryModel`.
- The same-file theorem `mobiusStrip_closedCylinder_boundary_models_not_homeomorphic` re-proves M-FINAL along that pipeline.

7.3 Instantiations

Theorems earlier in the paper are unconditional. This subsection records how the same statements bear on four debates outside pure mathematics. Where an extra premise is required, it is stated explicitly; nothing in the formalization defaults a controversial premise to true.

7.3.1 Consciousness and self-presence

Philosophical bridge (premise P2). Some accounts distinguish representational self-reference from a stronger notion of *self-presence* (first-person givenness, non-representational acquaintance, etc.). Premise **P2** is *not* a theorem in the library: it says only that, for the purposes of a topological comparison, the representational substrate is modeled by a space whose `ChartableR2`-detected boundary has one connectivity pattern (e.g. the cylinder's two disjoint components), while the target phenomenon is modeled by a space whose boundary has a different, incompatible pattern (e.g. the Möbius strip's single connected band—but the argument is *not specific to the Möbius strip*: Theorem 4.4 applies to *any* pair of spaces with incompatible `ChartableR2`-co-characterized boundaries, including Klein bottles, higher-genus non-orientable surfaces, or any other topological witness that satisfies the three hypotheses).

What is proved. Theorem 6.1 is unconditional: MobiusStrip $\not\simeq_t$ ClosedCylinder, as a concrete discharged witness. But the engine is Theorem 4.4: any two spaces whose `ChartableR2`-co-characterized boundaries are subtype-incompatible are not homeomorphic. Substituting a different target (Klein bottle, projective plane, etc.) for the Möbius strip does not evade the argument; it merely changes which boundary incompatibility is invoked. The adversary cannot escape by choosing a different non-orientable model.

Closed routes and admissibility. Multiple adversarial moves are blocked (see §7.6 for full detail):

- *Iterate deeper* (Horn 1): the regress is infinite, the Lawvere wall persists at every depth—more meta-levels do not close the gap.
- *Use a finite loop* (Horn 2): the iterate dichotomy shows bounded depth cannot host injective self-representation.
- *Represent via a different object* (Horn 3): section–retraction readout collapses to an idempotent on the host and the identity on the source; no injective tower on either side.
- *Use a partial self-model* (Horn 4): the diagonal blind spot means the information needed for self-capture is precisely the information the partial model cannot represent.
- *Restrict to decidable predicates* (Horn 6): the Bool Lawvere horn is closed; finite cardinality does not help.
- *Choose a different topological model* (Horn 5): the general lemma applies to any incompatible pair, not just cylinder/Möbius.

What remains **admissible** under P2: the relevant substrate must *already* have the target topological type (in which case it was never purely representational in the cylinder sense), or the account must invoke a discontinuous or non-topological bridge not modeled in this formalism. Scaling, iteration, richer self-models, or architectural reparameterization within a fixed topological type are *inadmissible* routes to the target.

7.3.2 Symbol grounding and semantic regress

In the symbol-grounding problem, arbitrary symbols must be interpreted, and interpreting the interpretation threatens a regress unless a non-representational stop is posited. Formally, a representational system already contains the morphism `represent` whose iterates index "one more representational step." The type alias `SymbolSystem` is definitionally identical to `RepresentationalSystem`; lemmas such as `symbolGrounding_metaRegressArrow_injective` record that distinct natural-number depths correspond to *distinct* endomorphisms of the awareness object.

Closed routes and admissibility. The horn closures constrain the symbol system on multiple fronts:

- If iterate indices are injective (Horn 1), interpretation levels form an infinite regress with the Lawvere wall intact at every stage.
- If iterates collide (Horn 2), the system has only finitely many distinct interpretation levels—a loop too shallow for full self-interpretation.
- If the system tries a *partial* interpretation function (Horn 4), the diagonal exclusion (`lawvere_diagonal_not_in_range`) guarantees a blind spot: the interpretation cannot represent its own diagonal behavior.
- Restricting to decidable predicates (Horn 6) does not help: the Bool-valued Lawvere horn is closed.

Admissible grounding must therefore come from *outside* the representational loop: the system cannot ground itself by iterating its own interpretive machinery, whether totally or partially, whether with infinite depth or finite.

7.3.3 Diagonal arguments and self-interpretation

Lawvere's theorem in `Type` implies that there is no family $s : A \to A \to B$ that realizes every unary $g : A \to B$ as $s(a)$ for some parameter a, once B carries an endomorphism without fixed points. `LawvereRegressBridge` collects the instances and connects them to the categorical presentation; `lawvere_fixed_point_stays_representational` says the `OntologicalSlot` wall persists at every iterate depth.

Closed routes and admissibility. The diagonal closures are cumulative:

- *Total self-interpreter* $s : A \to A \to B$: blocked by Lawvere for any B with a fixed-point-free endo (`no_universal_parametric_unary_nat` for $\mathbb{N}$, `no_universal_parametric_unary_bool` for Bool).
- *Partial self-interpreter*: abandoning surjectivity does not help. `lawvere_diagonal_not_in_range` shows that for *any* $s : A \to A \to B$—not necessarily surjective—the diagonal $a \mapsto f(s\,a\,a)$ is never $s\,a_0$ for any a_0. The blind spot is exactly the system's own diagonal behavior.
- *Restrict to finite / decidable predicates* (Horn 6): $\neg$: Bool $\to$ Bool is fixed-point-free, so the two-valued case is no easier.
- *Climb the tower*: however high the self-applicative coding rises, the endo $\neq$ obj distinction persists at every iterate depth.

Admissible self-interpreters must therefore be *incomplete*: they cannot model their own diagonal behavior. This is not a resource limitation (compute, training data, context length) but a structural one; it holds for every codomain B admitting a fixed-point-free endomorphism, and for every cardinality of A.

7.3.4 Quantum measurement and observer chains

Talk of apparatus, environment, and Wigner's-friend hierarchies suggests a chain of systems, each described relative to a larger "observer." If one ever gives such a chain a *topological* state space Obs, together with a subset ∂_{obs} whose membership agrees with $\neg$ ChartableR2, then Theorem 4.4 applies to any target space whose boundary subspace is not homeomorphic to ∂_{obs}—the Möbius strip is one such target, but the argument works for *any* space with incompatible boundary topology (Klein bottle, projective plane, higher non-orientable models, etc.). `QuantumObserverChainHypothesis` packages exactly this data as assumptions; the conclusion IsEmpty(Obs $\simeq_t$ Target) follows for any target satisfying the hypotheses.

Closed routes and admissibility. The same horn closures apply to the observer chain:

- Iterating the measurement chain (stacking more apparatus/environment layers) does not change the topological type of the state space (Horns 1–2).
- Encoding the observer inside a larger system via section–retraction does not produce an injective tower (Horn 3).
- Partial models of the observer–readout process still miss their own diagonal (Horn 4).
- Choosing a different non-orientable target does not evade Theorem 4.4 (Horn 5).

The chain must either introduce genuinely new topological structure at some level—not representational iteration of the same type—or accept that its boundary organization differs from the target.

The library does *not* define Hilbert spaces, observables, or projection postulates; any physical application requires an independent argument that a concrete model satisfies the hypothesis fields.

7.4 Scaling of language models and similar systems

Training larger models extends a representational architecture; it does not, by itself, produce a homeomorphism between a given topological model of that architecture and a different global model. Theorem 4.4 is indifferent to the choice of target: the cylinder/Möbius pair is one witness, but any target space with incompatible `ChartableR2`-detected boundary is equally blocked.

Closed routes and admissibility.

- More parameters, deeper layers, longer context windows: these extend the architecture *within* a fixed topological type (Horn 1).
- Finer-grained self-models built during training: partial self-interpreters still miss their own diagonal (Horn 4).
- Switching from infinite-depth to finite-depth loops (e.g. fixed-iteration transformers): the iterate dichotomy (Horn 2) shows bounded depth cannot host injective self-representation.
- Choosing a different target topology does not help: the general lemma covers any pair (Horn 5).

Admissible paths to a topological type with different boundary organization require *structural* change—a different topology of the underlying state space—not quantitative extension of the same architecture.

7.5 Relationship to the Trans-Computational Processing argument

An earlier line of work argues from computability—halting, Gödel-style diagonalization, and related limits—to constraints on "self-contained" computational mind models. The topological strand here is independent: it turns on `ChartableR2` sensors and boundary subspaces, not on recursiveness. The two approaches can reinforce a common *limitative* conclusion about representation closed under its own rules, but neither formally entails the other.

7.6 No-escape analysis: horn-by-horn closure

An adversary confronting the representational regress has, at first glance, several avenues of retreat. Each is formalized and shown to fail. The numbering below follows the order in which the escape was identified; each "horn" names an adversarial move, and the closing theorem blocks it.

Horn 1: Infinite injective tower. *Move:* accept iter_injective and try to use the tower itself as a self-representation resource. *Closure:* Theorem 2.2 yields infinitely many distinct meta-levels; Theorem 2.7 (the Lawvere wall) ensures the endo/obj distinction persists at every level. The tower does not collapse into closure.

Horn 2: Finite-order / periodic iterates. *Move:* weaken iter_injective—let $\text{represent}^n = \text{represent}^m$ for some $n \neq m$, so the loop has only finitely many distinct powers.
Closure: `IterateClosure.End.pow_iterate_dichotomy` proves the dichotomy: either the iterates are all distinct (Horn 1) or there exists a bound B such that every power equals some e^i with $i < B$. In the finite case the system's representational depth is bounded; it cannot host an injective tower at all.

Horn 3: Cross-object readout via section–retraction. *Move:* represent A *inside* a different object B using a section $s : A \to B$ and retraction $r : B \to A$ with $s \circ r = \mathrm{id}_A$. *Closure:* `CrossObjectRepresentation.section_retraction_no_injective_tower_either_side` shows that $r \circ s$ is idempotent on B (so $(r \circ s)^1 = (r \circ s)^2$, blocking injective powers on the host side) and $s \circ r = \mathrm{id}_A$ (so all powers are the identity on the source side—trivially periodic). Neither side supports an injective $\mathbb{N}$-tower.

Horn 4: Partial (non-surjective) self-model. *Move:* concede that no *total* self-interpreter exists, but claim a useful *partial* one. *Closure:* `lawvere_diagonal_not_in_range` requires no surjectivity of s: for *any* $s : A \to A \to B$ and any fixed-point-free $f : B \to B$, the diagonal function $a \mapsto f(s\,a\,a)$ is never in the range of s. Every partial self-model necessarily has a blind spot that includes its own diagonal—the information that matters most for self-capture is exactly the information it cannot represent.

Horn 5: Different topological models. *Move:* reject the cylinder/Möbius witness; use a different geometric model. *Closure:* Theorem 4.4 (`not_homeomorphic_of_chartableR2_boundary_contrast`) is *model-agnostic*: it applies to *any* pair of spaces whose ChartableR2-co-characterized boundaries are subtype-incompatible. The cylinder and Möbius strip are a discharged witness; any other pair meeting the same three hypotheses imports the same conclusion.

Horn 6: Computability / decidable self-reference. *Move:* restrict to decidable predicates or finite-valued functions, where "self-reference" seems more tractable than Cantorian diagonals. *Closure:* `no_universal_parametric_unary_bool` shows that no $s : A \to A \to \mathrm{Bool}$ can realize every $g : A \to \mathrm{Bool}$, since $\neg : \mathrm{Bool} \to \mathrm{Bool}$ is fixed-point-free. The obstruction is not an artifact of infinite cardinality; it holds in the decidable, two-valued setting.

Residual openings (not closed here). Two further routes are not formally addressed in this paper: *(i) higher-categorical self-maps* (a self-representation living in a 2-category or ∞-category, where the monoid of endomorphisms is replaced by a richer hom-object); and *(ii) non-categorical frameworks* (a system that resists categorical encoding entirely). The first is subject to analogous monoid-element arguments once a truncation to the 1-categorical endomorphism monoid is performed; the second falls outside the scope of any categorical formalization by definition. Both are noted as limitations in §8.

7.7 Admissibility: what the closures tell neighboring debates

Given that Horns 1–6 are closed, the formal landscape constrains what *admissible* solutions look like in each neighboring debate. The reasoning in every case is the same: the closures establish a structural barrier, so any proposed solution must either (a) operate outside the barrier's hypotheses, or (b) accept the consequences.

Consciousness and self-presence. If a system's state-space has cylinder-type boundary organization (two disjoint boundary components, each ¬ChartableR2), then no continuous global identification (homeomorphism) maps it to a Möbius-type target (one connected boundary band). An admissible account of self-presence must therefore either (i) posit that the relevant substrate *already* has the target topological type (in which case it was never purely representational in the cylinder sense), or (ii) invoke a discontinuous or non-topological bridge (which the formalism does not model and which would require independent justification). Scaling—more parameters, longer

context windows, deeper architectures—extends a representational system *within* a fixed topological type; it does not change boundary organization.

Symbol grounding. Any symbol system that iterates interpretation is caught in the Horns 1/2 dichotomy: either infinite distinct levels (full regress plus Lawvere wall) or finite-bounded depth (finite loop, no injective tower). Horn 4 adds that even a *partial* interpretation function misses its own diagonal. Admissible grounding must therefore come from *outside* the representational loop: the system cannot ground itself by iterating its own interpretive machinery, whether totally or partially.

Diagonal arguments / AI self-interpretation. Any proposed self-interpreter $s : A \to A \to B$ (total or partial) necessarily fails to include its own diagonal (`lawvere_diagonal_not_in_range`); this holds for $B = \mathbb{N}$ (numeric outputs), $B = \text{Bool}$ (decidable predicates), and any B with a fixed-point-free endomorphism. Admissible self-interpreters *must* be incomplete in a precise sense: they cannot model their own diagonal behavior. This is not a resource limitation but a structural one; no amount of compute, training data, or architectural ingenuity changes the diagonal exclusion.

Quantum measurement chains. If the measurement chain's state space has a boundary pattern matching the cylinder side of the development, then no homeomorphism identifies it with a Möbius-type target. The chain must either introduce genuinely new topological structure at some level—not representational iteration of the same type—or accept that its boundary organization differs from the target pattern.

LLM and AI scaling. Scaling extends a representational architecture within a fixed topological type. If that type has cylinder-side boundary organization, scaling does not produce a homeomorphism to a Möbius-type target. Admissible paths to a different topological type require *structural* change (different topology of the underlying state space), not quantitative extension (more of the same).

8 Limitations and scope

1. **Hypothesis, not derivation.** iter_injective is bundled, not proved from minimal data. Not every $f : A \to A$ in every category satisfies it. The regress is conditional on this non-triviality hypothesis.

2. **Lawvere in `Type`, not abstract CCC.** The formalization uses Type u and MonoidalClosed(Type u), not a general closed category over arbitrary C. A fully abstract CCC version remains future work.

3. **`OntologicalSlot` is definitional, not metaphysical.** The disjointness endo $\neq$ obj is trivially true by constructor discrimination. This is the point, but the triviality means the result is *semantic* (about the structure of type theory) rather than a deep theorem.

4. **Models are chosen, not unique.** The closed cylinder and Möbius strip are *specific* topological models chosen for clarity and tractability. A different philosophical account might use different geometric witnesses (e.g., Klein bottle, different quotient). The argument works for any two models separated by a topological invariant; these two are the cleanest.

5. **P2 is philosophical.** The premise "consciousness requires non-orientable topology" is not a theorem. Section 7.3 states only what follows when P2 is adjoined to the proved non-homeomorphism; accepting or rejecting P2 remains an independent issue.

6. **No** IsManifold **polish.** The Möbius strip is not given a full Mathlib IsManifold instance. The proof uses ChartableR2 (open neighborhoods homeomorphic to $\mathbb{R}^2$) as a self-contained predicate, sidestepping the need for a full manifold-with-boundary atlas. This is sufficient for the non-homeomorphism but less elegant than a full atlas approach.

7. **Route W is incomplete.** An alternative proof via circle-winding / doubling obstructions (Steps 1–4 in `CylinderMobiusNonhomeo.lean`) is partially formalized. Step 5 (transport under a hypothetical global homeomorphism) is not proved. The flagship proof via ChartableR2 is complete and does not depend on Route W.

8. `QuantumObserverChainHypothesis`**.** This definition only bundles topological assumptions (boundary set, `ChartableR2` biconditional,
subtype non-homeomorphism to the Möbius boundary). It does not import Hilbert spaces, operators, or measurement axioms; a physical application would need a separate argument that a concrete model satisfies those fields.

9. **Partial self-models: blind spot, not total impossibility.** `lawvere_diagonal_not_in_range` proves that every partial self-model misses its own diagonal; it does not prove that partial self-models are *useless*—only that they are necessarily incomplete in a precise, diagonal-shaped way. Whether that incompleteness matters for a given application is a separate question.

10. **Higher-categorical and non-categorical frameworks.** The horn closures rely on 1-categorical endomorphism monoids. A self-representation living in a 2-category or ∞-category replaces the monoid with a richer hom-object; the present arguments apply after truncation to the underlying 1-categorical monoid, but a native higher-categorical treatment is not given. Systems that resist categorical encoding entirely fall outside any categorical formalization by definition.

11. **Computability: decidable analogue only.** `no_universal_parametric_unary_bool` closes the Bool-valued Lawvere horn. A full treatment of Turing-degree obstructions (Rice's theorem, Kleene's recursion theorem) in the same framework is not attempted; the present result is the finite/decidable endpoint of the diagonal family.

9 Related work

9.1 Formalized topology

Mathlib's manifold and topology libraries underpin the half-space and chart lemmas; other proof assistants also host substantial topology. The weight here is on a concrete non-homeomorphism coupled to representational-system infrastructure, not on novel classification of differential invariants.

9.2 Lawvere's theorem in proof theory

Lawvere's fixed-point argument appears in multiple formal developments. `OntologicalSlot` and the wall lemmas supply complementary category-internal tagging alongside the `Type`-level diagonal

theorem.

9.3 Consciousness, Gödel, and AI

Gödel- and computability-based limits on mechanism (including Penrose-style arguments, discussions of strange loops and self-reference in the sense of Hofstadter, and information-theoretic theories of consciousness) occupy a large literature orthogonal to the cylinder/Möbius homeomorphism question, though informal exposition sometimes elides the difference.

9.4 Topology in philosophy of mind

Topological figures of speech appear in some philosophy of mind. The present work differs in that its central topological claims are proved and the generalization pattern is explicit.

10 Conclusion

The formalization delivers three interlocking layers:

(i) A machine-checked regress and Lawvere infrastructure (categorical strand),

(ii) A general ChartableR2 boundary-contrast lemma with a fully discharged compact witness MobiusStrip $\not\simeq_t$ ClosedCylinder (topological strand), and

(iii) A systematic no-escape analysis (Horns 1–6, §7.6) that formally closes the principal adversarial retreat routes: infinite regress, finite-order collapse, cross-object readout, partial self-models, model substitution, and decidable self-reference.

The central insight is not that infinite regress is bad—that is widely acknowledged. The insight is that the barrier is **topological and structural**, not computational or resource-bound. The obstacle is *wrong type*, not *inexhaustible queue*. A homeomorphism obstruction is a finite-data invariant: two spaces are either homeomorphic or not, regardless of how many modeling steps one enumerates. No amount of compute, no deeper self-model, no richer architecture changes the boundary pattern of a topological type. And the Lawvere diagonal exclusion (`lawvere_diagonal_not_in_range`) shows that even *partial* self-models necessarily have a blind spot containing their own diagonal—the very information needed for self-capture is the information they cannot represent.

What is not being claimed. The theorems are unconditional mathematics; the *application* to any particular domain requires domain-specific premises (P2 for consciousness, `ChartableR2` biconditionals for a specific state-space model, etc.). The formalism does not prove that consciousness is non-representational, that AI cannot be conscious, or that symbol grounding is impossible. It proves that *if* the relevant state space has the boundary pattern characterized here, *then* a homeomorphism to a space with a different boundary pattern does not exist—and that the standard adversarial attempts to work around that barrier fail in formally checkable ways.

The admissibility picture. The horn closures (§7.6) together with the admissibility analysis (§7.7) leave a clean landscape. In each neighboring debate—consciousness, symbol grounding, diagonal limitations, measurement chains, AI scaling—the admissible solutions share a common shape: *structural change* to the topological type of the substrate, not *quantitative extension* within

a fixed type. The formalism makes this precise and machine-checkable, so that future claims about circumventing the barrier can be evaluated against a fixed standard.

Future work. A full IsManifold instance for MobiusStrip; completion of the optional Route W winding track; a more abstract Lawvere statement for general cartesian closed categories; native higher-categorical horn closures (2-category and ∞-category analogues); further models where ChartableR2 boundary data can be established; and a full Turing-degree analogue of the decidable Lawvere horn.

A Lean Module Map

Module	**Role**
`Basic`	`RepresentationalSystem`, iterates, `OntologicalSlot`
`Regress`	Non-termination, unbounded levels
`FixedPoints`	`EndoVsPoint`, Lawvere wall non-dissolution
`Lawvere`	Fixed-point theorem in `Type`
`LawvereCCCType`	`MonoidalClosed` packaging
`Orientability`	T_2 homeomorphism invariance
`CylinderMobius`	Fundamental domain, quotient, compact Hausdorff, boundary band
`CylinderBoundary`	Closed cylinder boundary: two disjoint connected faces
`CylinderChartable Boundary`	Cylinder C4 biconditional
`ChartableR2Neighbor`	`ChartableR2`, `ChartableR2BoundaryModel`; openness / stability
`ChartableR2Bridge`	`Homeomorph.chartableR2_iff`; `not_homeomorphic_of_chartableR2_boundary_contrast`; `ChartableR2BoundaryModel.not_homeomorphic`
`ChartableR2Concrete BoundaryModels`	Möbius/cylinder boundary models; conditional observer chain; `Quantum ObserverChainHypothesis`
`SymbolGrounding`	`SymbolSystem`; `symbolGrounding_*`
`LawvereRegressBridge`	`no_universal_parametric_unary_nat`, `no_universal_parametric_unary_bool`; `lawvere_diagonal_not_in_range`; `lawvere_fixed_point_stays_representational`
`NoEscapeClosure`	Iterate collisions vs. injective towers; section–retraction idempotent readout
`ChartableR2Models`	Interior charts; cylinder interior `ChartableR2`
`MobiusChartable Boundary`	Open-cell charts; boundary ¬`ChartableR2`; Möbius C4 packaging
`MobiusSeamChartable`	Saturated seam patches (open, saturated)
`MobiusSeamChart`	IFT-based local model; trig coordinates
`MobiusSeamTrigInject`	Trig injectivity on quotient windows
`MobiusSeamChartableR2`	Seam charts (trig + unfold); interior-height; **M-FINAL**
`MobiusCylinder BoundaryContrast`	Boundary subspaces not homeomorphic
`CylinderMobius Nonhomeo`	Route W (optional; Steps 1–4)
`HalfLineVsLine`	1D half-space $\neq$ line
`HalfPlaneVsPlane`	2D half-space $\neq$ plane
`HalfSpaceNeighborVs Plane`	Local half-space $\neq$ open plane patch
`PuncturedPlaneNot SimplyConnected`	$\mathbb{C} \setminus \{0\}$ not simply connected
`OpenCompactChart Obstruction`	Compact $\mathbb{R}^2$ patch $\neq \mathbb{R}^2$
`Main`	Master theorem + extended bundle

B Key Theorem Index

C Reproduction

```
cd representational-regress-lean
lake update
lake exe cache get
lake build RepresentationalRegress
# Verify: 0 sorry
rg "sorry" RepresentationalRegress/
```

Toolchain: `leanprover/lean4:v4.29.0-rc6`. Mathlib: `v4.29.0-rc6`.

Lean name	Statement	Module
regress_iterates_unbounded	$\forall B, \exists L > B$, distinct iterate	Regress
lawvere_fixed_point_Type	Universal enumerator $\Rightarrow$ every endo has fixed point	Lawvere
lawvere_wall_is_not_dissolution	endo(represent) $\neq$ obj(fp)	FixedPoints
instT2SpaceMobiusStrip	Möbius strip is T_2	CylinderMobius
isConnected_mobiusStripBoundary	Boundary band is connected	CylinderMobius
closedCylinder_boundary_eq_union	Cylinder boundary = 2 disjoint faces	CylinderBoundary
euclideanHalfSpace2_not_homeomorphic_euclidean2	$H^2 \not\cong_t \mathbb{R}^2$	HalfPlaneVsPlane
not_homeomorphic_of_chartableR2_boundary_contrast	Generic: incompatible boundaries $\Rightarrow \neg$ homeo	ChartableR2Bridge
closedCylinder_boundaryUnion_iff_not_chartableR2	Cylinder C4 biconditional	CylinderChartableBoundary
not_chartableR2_of_mem_mobiusStripBoundary	Möbius boundary $\Rightarrow \neg$ChartableR2	MobiusChartableBdy
chartableR2_mobiusQuotientMk_of_interior_height	Interior height $\Rightarrow$ ChartableR2	MobiusSeamChartableR2
mobiusStripBoundary_not_homeomorphic...	Boundary subtypes not homeo	MobiusCylBdyContrast
mobiusStrip_not_homeomorphic_closedCylinder	**M-FINAL**	MobiusSeamChartableR2
representational_regress_master	Core master conjunction	Main
representational_regress_master_extended	Extended bundle (+ topology)	Main
QuantumObserverChainHypothesis.not_homeomorphic_mobiusStrip	Conditional: obs. chain $\not\cong_t$ Möbius	ChartableR2ConcreteBoundaryModels
no_universal_parametric_unary_nat	No universal $s : A \to A \to \mathbb{N}$ for all $g : A \to \mathbb{N}$	LawvereRegressBridge
IterateClosure.End.pow_iterate_dichotomy	Injective iterate powers vs. finite index bound	NoEscapeClosure
CrossObjectRepresentation.section_retraction_no_injective_tower_either_side	Both-side: $\neg$ injective $\mathbb{N}$-powers on source *or* host	NoEscapeClosure
lawvere_diagonal_not_in_range	Diagonal never in range of any $s : A \to A \to B$ (no surjectivity)	LawvereRegressBridge
no_universal_parametric_unary_bool	No universal $s : A \to A \to \mathrm{Bool}$ for all $g : A \to \mathrm{Bool}$	LawvereRegressBridge